Implementation of static and dynamic semantics for a calculus with algebraic effects and handlers using `PLT Redex`

(Implementacja statycznej i dynamicznej semantyki rachunku z efektami algebraicznymi i ich obsługą z pomocą biblioteki `PLT Redex`)

Maciej Buszka

Praca inżynierska

Promotor: dr hab. Dariusz Biernacki

Abstract

English abstract

Abstrakt w języku polskim

Contents

Chapter 1

Introduction

Algebraic effects [6] are an increasingly popular technique of structuring computational effects. They allow for seamless composition of multiple effects, while retaining (unlike monads) applicative style of programs. Coupled with handlers [7] which give programmers ability to interpret effects, they provide a great tool for abstracting over a set of operations which some program may perform and separating this interface from semantics of those operations defined as effect handlers.

As these features are highly desirable many calculi and languages have been developed in order to get them just right; most notable of them being: *Koka* (featuring type inference, effect polymorphism with row-types and *Javascript*-like syntax), *Links* (featuring *ML*-like syntax, row-typed effect polymorphism and ad-hoc effects) and *Eff* [1] with implicit effect checking and recent work on direct compilation to *OCaml* [4]. On more theoretical side, various approaches to semantics of algebraic effects can be spotted in literature, both in respect to type system and run-time semantics. Although most calculi use some form of row-types to implement tracking of effects there are differences in permitted shapes (at most one effect of given type or many effects), whether effects must be defined before use or not and how effects interact with polymorphism and abstraction. At run-time handlers can wrap the captured continuation (giving so-called deep handlers) or not (shallow handlers) and the very act of finding the right handler can be implemented in various ways, mainly depending on some constructs which skip handlers.

All this variety naturally invites Us to experiment with different features and components of a calculus. In this thesis I will build such a calculus, describing and justifying my choices and discussing the trade-offs I faced. In order to to rapidly iterate on design and test the calculus, I decided to use `PLT Redex` library which allows for building executable type system judgments and reduction relation. As such the other goal of my thesis was to assess viability of `PLT Redex` for development of bigger calculi. To briefly summarize my development, the calculus consists of:

- Curry style type system with ad-hoc effects in the style of *Links*, effect rows

based on *Koka* and *lift* construct first shown in [2], implemented as unification
based type inference algorithm.

- Executable reduction semantics most similar to system of [2].

- CEK style abstract machine with stack and *meta*-stack of handlers

- Language front-end which translates human-friendly programs to calculus terms,
  integrated with `Racket` environment, which allows for easy experimentation.

The rest of this thesis is structured as follows: In chapter 2 I describe the
calculus in greater detail, in chapter 3 I discuss technicalities of implementation,
in chapter 4 I summarize the process of integration with `Racket` environment and
chapter 5 is user's manual. In the reminder of this chapter I introduce main topics
of this thesis.

1.1   Algebraic effects and handlers

Algebraic effects and handlers are a language level framework which allow for co-
herent presentation, abstraction, composition and reasoning about computational
effects. The key idea is to separate invocation of an effectful operation in an expres-
sion from the meaning of such operation. When one invokes an operation, current
continuation (up to nearest handler) is captured and passed along with operation's
argument to nearest handler. The handler in turn may execute arbitrary expression,
using the continuation once, twice, returning a function which calls the continuation
or simply ignoring it. This way many control structures can be modeled and gen-
eralized by algebraic effects and appropriate handlers. For example, the exceptions
can be modeled using a single operation `throw` and a handler which either returns
the result when computation succeeded or returns default value, ignoring passed
continuation.

```
handle e with
| throw () r -> // return default value
| return x   -> x
end
```

From the language design standpoint algebraic effects provide single implementation
of various phenomena which may happen during execution of a program, for example
mutable state, I/O, environment lookup, exceptions etc. in a sense that every effect
is treated the same, the typing rules are defined for invocation of any operation,
and handling of any operation. Similarly the operational semantics is also quite
simple and succinct thanks to uniform treatment of various effects. This framework
is also extendable. With small additions if can handle built-in effects in addition to
user-defined ones.

From the language user perspective algebraic effects provide means of abstraction over effects used in a program. Thanks to easy creation of new effects, one can define special purpose operations and their handlers to better represent domain specific problems while simultaneously using well known effects, defined in standard library. With effects being tracked by the type system, programmers can enforce purity or specific set of used effects at compile-time, or using effect polymorphism they can write reusable functions which abstract over effects which may happen. The separation of definition and implementation of effects allows for various interpretations of operations, similar to a technique of *dependency-injection* used for example during testing.

1.2 Type inference

Type inference is a technique of algorithmic reconstruction of types for various constructions used in a language. It allows programmers to write programs with no type annotations, which often feel redundant and obfuscate the meaning of a program. The most well known type system with inference is a system for *ML* family of languages - *Haskell*, *OCaml*, *SML* which infers the types with no annotations whatsoever. Formal type system defines grammar of types consisting of base types (`int`, `bool` etc.), type constructors (arrows, algebraic data types) and type variables. The typing rules require types which should be compatible (f.e. formal parameter and argument types) to unify. The key feature of this system is so called let-polymorphism - generalization of types of let-bound variables. This way code reuse can be accomplished without complicating the type system and compromising type safety. The basis of implementation of this system is first order unification algorithm, which syntactically decomposes types and builds a substitution from type variables to types.

1.3 Reduction semantics and abstract machines

Reduction semantics is a format for specifying dynamic semantics of a calculus in an operational style. The basic idea is to first define redexes - expressions which can be reduced, and contexts in which the reduction can happen.

$$e ::= x \mid \lambda x.e \mid ee$$
$$E ::= \Box \mid Ee$$

Then we have to define reduction relation, which generates a transition system for terms in the calculus. The rules usually take the form of redex in a context which reduces to some expression in a context.

1.4 PLT Redex

Chapter 2

The calculus

The calculus implemented in this thesis is based on lambda calculus with call-by-value semantics, similarly to other calculi which allow for computational effects, because fixed evaluation order is essential to obtaining sane program semantics. Inspired by Links [3] the operations are truly ad-hoc meaning that they don't have to be declared before usage. Moreover the calculus requires no type annotations whatsoever in spirit of *ML* family of languages while still tracking effects which occur in a program.

2.1 Abstract syntax

Abstract syntax of comprises of forms standard to λ-calculus (λ abstractions, variables and applications), extended with number literals and primitive numerical operations (which allow for some example computation), and syntactic forms used by algebraic effects - operation invocations, handle expressions and lift expressions. Inspired by calculus of Links described in ??? operations are a distinct syntactic category, while lift construct was introduced in Biernacki et. al.

2.2 Static semantics

The type system is based on Koka (Leijen's style of row types) and Links systems (ad-hoc operations). As the calculus is defined in Curry-style the inference rules describe type reconstruction algorithm.

2.2.1 Row types

Row types as described in [5]

$$
\begin{aligned}
v &::= number \mid (\lambda\ x\ e) \\
prim &::= +\mid -\mid * \\
e &::= v \mid x \mid (e\ e) \\
&\qquad \mid (op\ e) \mid (\text{handle}\ e\ hs\ ret) \\
&\qquad \mid (\text{lift}\ op\ e) \mid (prim\ e\ e) \\
hs &::= ((op_{!\_1}\ hexpr)\ ...) \\
hexpr &::= (x_{!\_1}\ x_{!\_1}\ e) \\
ret &::= (\text{return}\ x\ e) \\
h &::= (op\ hexpr) \\
t &::= \text{Int} \mid (t \text{ -> } row\ t) \mid (t \text{ => } t) \mid row \mid a \\
row &::= (op\ t\ row) \mid a \mid \cdot \\
x &::= (\text{variable-prefix v:}) \\
a &::= (\text{variable-prefix t:}) \\
op &::= (\text{variable-prefix op:}) \\
\Gamma &::= (x\ t\ \Gamma) \mid \cdot \\
E &::= [] \mid (E\ e) \mid (v\ E) \mid (prim\ E\ e) \mid (prim\ v\ E) \\
&\qquad \mid (op\ E) \mid (\text{handle}\ E\ hs\ ret) \\
&\qquad \mid (\text{lift}\ op\ E) \\
S &::= (a\ t\ S) \mid \cdot \\
N, n &::= natural \\
SN &::= (S\ N)
\end{aligned}
$$

2.2.2 Type inference

The main judgment *infer* infers a type and effect row, and calculates new substitution, given typing environment, current substitution and an expression. As in *ML* languages only simple types can be inferred, along with effect rows

2.2.3 Effect handlers

2.3 Dynamic semantics

2.4 Abstract machine

Chapter 3

Implementation

3.1 `PLT Redex`

3.2 Typing relation

3.3 Unification

3.4 Reduction relation

3.5 Automatic testing

Chapter 4

The `Racket` environment

4.1   Front-end

4.2   Back-end

Chapter 5

User's manual

Bibliography

[1] Andrej Bauer and Matija Pretnar. ?Programming with Algebraic Effects and Handlers? In: *CoRR* abs/1203.1539 (2012). arXiv: 1203.1539. URL: http://arxiv.org/abs/1203.1539.

[2] Dariusz Biernacki et al. ?Handle with Care: Relational Interpretation of Algebraic Effects and Handlers? In: *Proc. ACM Program. Lang.* 2.POPL (Dec. 2017), 8:1–8:30. ISSN: 2475-1421. DOI: 10.1145/3158096. URL: http://doi.acm.org/10.1145/3158096.

[3] Daniel Hillerström and Sam Lindley. ?Liberating Effects with Rows and Handlers? In: *Proceedings of the 1st International Workshop on Type-Driven Development.* TyDe 2016. Nara, Japan: ACM, 2016, pp. 15–27. ISBN: 978-1-4503-4435-7. DOI: 10.1145/2976022.2976033. URL: http://doi.acm.org/10.1145/2976022.2976033.

[4] Oleg Kiselyov and KC Sivaramakrishnan. ?Eff Directly in OCaml? In: *arXiv e-prints*, arXiv:1812.11664 (Dec. 2018), arXiv:1812.11664. arXiv: 1812.11664 [cs.PL].

[5] Daan Leijen. ?Extensible records with scoped labels? In: *Proceedings of the 2005 Symposium on Trends in Functional Programming (TFP'05), Tallin, Estonia.* Sept. 2005. URL: https://www.microsoft.com/en-us/research/publication/extensible-records-with-scoped-labels/.

[6] Gordon Plotkin and John Power. ?Algebraic Operations and Generic Effects? In: *Applied Categorical Structures* 11.1 (Feb. 2003), pp. 69–94. ISSN: 1572-9095. DOI: 10.1023/A:1023064908962. URL: https://doi.org/10.1023/A:1023064908962.

[7] Gordon Plotkin and Matija Pretnar. ?Handling Algebraic Effects? In: *Logical Methods in Computer Science* 9.4 (Dec. 2013). Ed. by Andrzej Tarlecki. ISSN: 1860-5974. DOI: 10.2168/lmcs-9(4:23)2013. URL: http://dx.doi.org/10.2168/LMCS-9(4:23)2013.