Implementation of static and dynamic semantics for a calculus with algebraic effects and handlers using `PLT Redex`

(Implementacja statycznej i dynamicznej semantyki rachunku z efektami algebraicznymi i ich obsługą z pomocą biblioteki `PLT Redex`)

Maciej Buszka

Praca inżynierska

Promotor: dr hab. Dariusz Biernacki

Abstract

English abstract

_____

Abstrakt w języku polskim

Contents

Chapter 1

Introduction

Algebraic effects [8] are an increasingly popular technique of structuring computational effects. They allow for seamless composition of multiple effects, while retaining (unlike monads) applicative style of programs. Coupled with handlers [9] which give programmers ability to interpret effects, they provide a great tool for abstracting over a set of operations which some program may perform and separating this interface from semantics of those operations defined as effect handlers.

As these features are highly desirable many calculi and languages have been developed in order to get them just right; most notable of them being: *Koka*[7] (featuring type inference, effect polymorphism with row-types and *Javascript*-like syntax), *Links*[4] (featuring *ML*-like syntax, row-typed effect polymorphism and ad-hoc effects) and *Eff*[1] with implicit effect checking and recent work on direct compilation to *OCaml*[5]. On more theoretical side, various approaches to semantics of algebraic effects can be spotted in literature, both in respect to type system and run-time semantics. Although most calculi use some form of row-types to implement tracking of effects there are differences in permitted shapes (at most one effect of given type or many effects), whether effects must be defined before use or not and how effects interact with polymorphism and abstraction. At run-time handlers can wrap the captured continuation (giving so-called deep handlers) or not (shallow handlers) and the very act of finding the right handler can be implemented in various ways, mainly depending on some constructs which skip handlers.

All this variety naturally invites Us to experiment with different features and components of a calculus. In this thesis I will build such a calculus, describing and justifying my choices and discussing the trade-offs I faced. In order to rapidly iterate on design and test the calculus, I decided to use `PLT Redex` library which allows for building language model with executable type system judgments and reduction relation. As such the other goal of my thesis was to assess viability of `PLT Redex` for development of bigger calculi. To briefly summarize my development, the calculus consists of:

- Curry style type system with ad-hoc effects in the style of *Links*, effect rows based on *Koka* and *lift* construct first shown in [2], implemented as unification based type inference algorithm.

- Executable reduction semantics most similar to system of [2].

- CEK style abstract machine with stack and *meta*-stack of handlers

- Language front-end which translates human-friendly programs to calculus terms, integrated with `Racket` environment, which allows for easy experimentation.

The rest of this thesis is structured as follows: in chapter 2 I describe the calculus in greater detail, in chapter 3 I discuss technicalities of implementation, in chapter 4 I summarize the process of integration with `Racket` environment and chapter 5 is user's manual. In the reminder of this chapter I introduce main topics of this thesis.

1.1   Algebraic effects and handlers

Algebraic effects and handlers are a language level framework which allow for coherent presentation, abstraction, composition and reasoning about computational effects. The key idea is to separate invocation of an effectful operation in an expression from the meaning of such operation. When one invokes an operation, current continuation (up to nearest handler) is captured and passed along with operation's argument to nearest handler. The handler in turn may execute arbitrary expression, using the continuation once, twice, returning a function which calls the continuation or simply ignoring it. This way many control structures can be modeled and generalized by algebraic effects and appropriate handlers. For example, the exceptions can be modeled using a single operation `Throw` and a handler which either returns the result when computation succeeded or returns default value, ignoring passed continuation.

```
handle e with
| Throw () r -> // return default value
| return x   -> x
end
```

From the language design standpoint algebraic effects provide single implementation of various phenomena which may happen during execution of a program, for example mutable state, I/O, environment lookup, exceptions etc. in a sense that every effect is treated the same, the typing rules are defined for invocation of any operation, and handling of any operation. Similarly the operational semantics is also quite simple and succinct thanks to uniform treatment of various effects. This framework

is also extendable. With small additions if can handle built-in effects in addition to user-defined ones.

From the language user perspective algebraic effects provide means of abstraction over effects used in a program. Thanks to easy creation of new effects, one can define special purpose operations and their handlers to better represent domain specific problems while simultaneously using well known effects, defined in standard library. With effects being tracked by the type system, programmers can enforce purity or specific set of used effects at compile-time, or using effect polymorphism they can write reusable functions which abstract over effects which may happen. The separation of definition and implementation of effects allows for various interpretations of operations, similar to a technique of *dependency-injection* used for example during testing.

1.2 Type inference

Type inference is a technique of algorithmic reconstruction of types for various constructions used in a language. It allows programmers to write programs with no type annotations, which often feel redundant and obfuscate the meaning of a program. The most well known type system with inference is a system for *ML* family of languages - *Haskell*, *OCaml*, *SML* which infers the types with no annotations whatsoever. Formal type system defines grammar of types consisting of base types (`int`, `bool` etc.), type constructors (arrows, algebraic data types) and type variables. The typing rules require types which should be compatible (f.e. formal parameter and argument types) to unify. The key feature of this system is so called let-polymorphism - generalization of types of let-bound variables. This way code reuse can be accomplished without complicating the type system and compromising type safety. The basis of implementation of this system is first order unification algorithm, which syntactically decomposes types and builds a substitution from type variables to types.

1.3 Reduction semantics and abstract machines

Reduction semantics is a format for specifying dynamic semantics of a calculus in an operational style. The basic idea is to first define redexes - expressions which can be reduced, and contexts in which the reduction can happen. Taking λ-calculus with call-by-value reduction order as an example, the only redex is application of a function to value $(\lambda x.e)v$. The possible contexts are: empty context $\square$ or evaluation of operator part of application Ee or evaluation of operand vE. With these possibilities in mind, we will define binary relation $\longrightarrow$ which describes single step of reduction. Such relation can be thought of as a transition system, rewriting terms into simpler ones step by step. There usually are two approaches to definition of such relation:

v ::= (λ *x* *e*) | *number*
e ::= *v* | (*e* *e*) | *x*
x ::= *variable-not-otherwise-mentioned*
E ::= [] | (*E* *e*) | (*v* *E*)

Figure 1.1: λ-calculus abstract syntax

E[((λ *x* *e*) *v*)] ⟶ *E*[substitute⟦*e*, *x*, *v*⟧] [β]

Figure 1.2: λ-calculus reduction relation

- Definition of primitive reduction $(\lambda x.e)v \longrightarrow_p e\{v/x\}$ which operates only on redexes and giving it a closure via following inference rule: $\dfrac{e \longrightarrow_p e'}{E[e] \longrightarrow E[e']}$, which says that if we can primitively reduce some expression, than we can do it in any context.

- Or definition of $\longrightarrow$ directly, with decomposition of terms on both sides: $E[(\lambda x.e)v] \longrightarrow E[e\{v/x\}]$

where the syntax $e\{v/x\}$ means term e with value v substituted for variable x, and $E[e]$ means some context E with expression e inserted into the hole. For both approaches it is important, that any term can be uniquely decomposed into redex and context, because when it is the case, then the relation is deterministic and gives good basis for formulation of abstract machines, interpreters or transformations to some other intermediate representations.

Abstract machine is a mathematical construction, usually defined as a set of configurations with deterministic transformations, which are computationally simple. The goal for formulation of an abstract machine is to mechanize evaluation of terms while retaining semantics given in more abstract format, f.e. reduction semantics, with the correspondence being provable[3]. As an example I will show a *CEK*-machine for the λ-calculus defined earlier. The name *CEK* comes from *C*ommand, *E*nvironment and *K*ontinuation. The machine configuration is a triple (e, ρ, κ) where e is an expression which is decomposed or reduced, ρ is an environment mapping variables to values, and the last component κ is a continuation stack, which determines what will happen with value, to which first component eventually reduces. Thanks to the environment we no longer have to explicitly perform substitution, leading to more machine friendly and efficient implementation. Given an initial state, the machine can then repeatedly apply transformation relation, either looping, arriving at a final value, or getting stuck.

1.4   PLT Redex

λ-calculus reduction example

Figure 1.3: λ-calculus example reduction sequence

CEK-machine for λ-calculus

Figure 1.4: λ-calculus abstract machine

Chapter 2

The calculus

The calculus implemented in this thesis is based on lambda calculus with call-by-value semantics. This choice follows other calculi which allow for computational effects, because fixed evaluation order is essential to obtaining sane program semantics. Inspired by *Links* [4] the operations are truly ad-hoc meaning that they don't have to be declared before usage. Moreover the calculus requires no type annotations whatsoever in spirit of *ML* family of languages while still tracking effects which occur in a program. The λ-calculus is extended with base types (Numeric and Boolean) with corresponding operations and literals, conditional expression and recursive functions. Effectful operations are invoked similarly to function calls (although they are a distinct syntactic category) and are handled using *handle* construct, they may also be lifted with *lift* syntactic form. Usually the operation will be handled by the closest handler with appropriate case, unless there is a lift for this operation in between, each one causing operation to skip one handler. Abstract syntax of the calculus is presented in Figure 2.1

2.1 Static semantics

The type system is based on *Koka*(Leijen's style of row types[6]), *Links*[4] (ad-hoc operations) and Biernacki et al. [2] (lift construct) systems. Initially I implemented a variant of System-F extended with row-types but it proved to be a bit of a mouthful to write even simplest programs. Moreover the `Redex`'s facilities for automatic testing were not able to generate well typed terms, so some type inference was inevitable. To limit the amount of work I decided to present the calculus in Curry style, with typing relation inferring the type for unannotated terms. Building on well known foundations, types are inferred via first order unification. The system does not feature polymorphism in first class fashion, as there is no rule where types are generalized, but I believe it to be a straightforward addition, following the *Koka*[7] calculus. Still, after inferring type of an expression, we can see which unification variables are left abstract and could be generalized. There are two main features

```
        b ::= true | false
        m ::= number
        v ::= m | b | (λ x e) | (rec x x e)
     prim ::= + | - | * | == | <= | >=
        e ::= v | x | (e e) | (if e e e) | (fix e)
              | (op e) | (handle e hs ret) | (lift op e) | (prim e ...)
       hs ::= ((op!_1 hexpr) ...)
    hexpr ::= (x!_1 x!_1 e)
      ret ::= (return x e)
        h ::= (op hexpr)
        t ::= Int | Bool | (t -> row t) | (t => t) | row | a
      row ::= (op t row) | a | ·
        x ::= (variable-prefix v:)
        a ::= (variable-prefix t:)
       op ::= (variable-prefix op:)
        Γ ::= (x t Γ) | ·
        E ::= [] | (E e) | (v E) | (prim E e) | (prim v E) | (if E e e)
              | (op E) | (handle E hs ret) | (lift op E)
        S ::= (a t S) | ·
     N, n ::= natural
       SN ::= (S N)
```

Figure 2.1: Abstract syntax

differentiating this system from *Koka*'s; firstly effects need not be defined before use, their signature is inferred as with any other construction; secondly the system is algorithmic, with rules explicitly encoding a recursive function which can infer the type of an expression.

2.1.1 Row-types

Row types as described in [6]

2.1.2 Type inference

The judgment $\Gamma \mid [S_1 \, N_1] \vdash e : t \, ! \, row \mid [S_2 \, N_2]$ asserts that in typing context Γ under type substitution S_1, with name supply state N_1 expression e has type t with effects *row* under type substitution S_2 and with name supply state S_2. Algorithmically this judgment infers a type and effect row, and calculates new substitution, given typing environment, current substitution and an expression. As in *ML* languages only simple types can be inferred, along with effect rows.

    Base rules for constants and variable lookup are straightforward, each introducing fresh effect row variable. To check λ expression, we first introduce fresh type variable, and then check the body in extended environment. The arrow gets annotated with effects which may occur during evaluation of the body and the λ abstraction itself is returned with fresh effect row. The recursive function checking is similar to normal functions. First variable denotes function itself, while second

$$\frac{\text{fresh-row}[\![N_1, row, N_2]\!]}{\Gamma \mid [S\ N_1] \vdash b : \text{Bool} \mathbin{!} row \mid [S\ N_2]}$$

$$\frac{\text{fresh-row}[\![N_1, row, N_2]\!]}{\Gamma \mid [S\ N_1] \vdash m : \text{Int} \mathbin{!} row \mid [S\ N_2]}$$

$$\frac{\text{lookup}[\![\Gamma, x, t]\!] \quad \text{fresh-row}[\![N_1, row, N_2]\!]}{\Gamma \mid [S\ N_1] \vdash x : t \mathbin{!} row \mid [S\ N_2]}$$

$$\frac{\text{fresh-var}[\![N_1, t_1, N_2]\!] \quad \text{fresh-row}[\![N_2, row_2, N_3]\!]}{(x\ t_1\ \Gamma) \mid [S_1\ N_3] \vdash e : t_2 \mathbin{!} row_1 \mid SN}{\Gamma \mid [S_1\ N_1] \vdash (\lambda\ x\ e) : (t_1 \to row_1\ t_2) \mathbin{!} row_2 \mid SN}$$

$$\frac{\text{fresh-arr}[\![N_1, t_1, \to, row_1, t_2, N_2]\!] \quad \text{fresh-row}[\![N_2, row_2, N_3]\!] \\ (x_f\ (t_1 \to row_1\ t_2)\ (x_a\ t_1\ \Gamma)) \mid [S_1\ N_3] \vdash e : t \mathbin{!} row \mid SN_1 \\ SN_1\ row_1 \sim row\ SN_2 \quad SN_2\ t_2 \sim t\ SN_3}{\Gamma \mid [S_1\ N_1] \vdash (\text{rec}\ x_f\ x_a\ e) : (t_1 \to row_1\ t_2) \mathbin{!} row_2 \mid SN_3}$$

$$\frac{\Gamma \mid SN_1 \vdash e_1 : t_a \mathbin{!} row_a \mid SN_2 \quad \text{unify-arr}[\![SN_2, t_a, t_1, \to, row_1, t_2, SN_3]\!] \\ \Gamma \mid SN_3 \vdash e_2 : t_3 \mathbin{!} row_2 \mid SN_4 \\ SN_4\ t_1 \sim t_3\ SN_5 \quad SN_5\ row_1 \sim row_2\ SN_6 \quad SN_6\ row_1 \sim row_a\ SN_7}{\Gamma \mid SN_1 \vdash (e_1\ e_2) : t_2 \mathbin{!} row_2 \mid SN_7}$$

$$\frac{\text{check-prim}[\![\Gamma, SN_1, prim, (e\ ...), t, row, SN_2]\!]}{\Gamma \mid SN_1 \vdash (prim\ e\ ...) : t \mathbin{!} row \mid SN_2}$$

$$\frac{\Gamma \mid SN_1 \vdash e_{cond} : t_{cond} \mathbin{!} row_{cond} \mid SN_2 \quad SN_2\ t_{cond} \sim \text{Bool}\ SN_3 \\ \Gamma \mid SN_3 \vdash e_{then} : t_{then} \mathbin{!} row_{then} \mid SN_4 \quad \Gamma \mid SN_4 \vdash e_{else} : t_{else} \mathbin{!} row_{else} \mid SN_5 \\ SN_5\ t_{then} \sim t_{else}\ SN_6 \quad SN_6\ row_{cond} \sim row_{then}\ SN_7 \quad SN_7\ row_{then} \sim row_{else}\ SN_8}{\Gamma \mid SN_1 \vdash (\text{if}\ e_{cond}\ e_{then}\ e_{else}) : t_{then} \mathbin{!} row_{then} \mid SN_8}$$

$$\frac{\Gamma \mid SN_1 \vdash e : t_1 \mathbin{!} row_1 \mid [S_1\ N_1] \quad \text{fresh-row}[\![N_1, row_2, N_2]\!] \quad \text{fresh-var}[\![N_2, t_2, N_3]\!] \\ [S_1\ N_3]\ (op\ (t_1 \Rightarrow t_2)\ row_2) \sim row_1\ SN_2}{\Gamma \mid SN_1 \vdash (op\ e) : t_2 \mathbin{!} row_1 \mid SN_2}$$

$$\frac{\Gamma \mid SN_1 \vdash e : t \mathbin{!} row \mid [S_1\ N_1] \quad \text{fresh-var}[\![N_1, a, N_2]\!]}{\Gamma \mid SN_1 \vdash (\text{lift}\ op\ e) : t \mathbin{!} (op\ a\ row) \mid [S_1\ N_2]}$$

$$\frac{\Gamma \mid SN_1 \vdash e : t_1 \mathbin{!} row_1 \mid SN_2 \quad (x\ t_1\ \Gamma) \mid SN_2 \vdash e_{ret} : t_{ret} \mathbin{!} row_{ret} \mid SN_3 \\ \text{infer-handlers}[\![\Gamma, SN_3, t_{ret}, hs, row_{out}, row_{handled}, SN_4]\!] \quad SN_4\ row_{out} \sim row_{ret}\ SN_5 \\ SN_5\ row_1 \sim row_{handled}\ SN_6}{\Gamma \mid SN_1 \vdash (\text{handle}\ e\ hs\ (\text{return}\ x\ e_{ret})) : t_{ret} \mathbin{!} row_{out} \mid SN_6}$$

Figure 2.2: Type system

$$\frac{\begin{array}{c} \text{fresh-row}[\![N_1,\ a_{handled},\ N_3]\!] \end{array}}{\text{infer-handlers}[\![\Gamma,\ [S\ N_1],\ t,\ (),\ a_{handled},\ a_{handled},\ [S\ N_3]]\!]}$$

$$\frac{\begin{array}{c} \text{trace}[\![(\text{"=== inferring handlers ===}\backslash n{\sim}s\backslash n\text{"}\ ((op\ (x_v\ x_r\ e))\ h\ ...))]\!] \\ \text{infer-handlers}[\![\Gamma,\ SN_1,\ t_{ret},\ (h\ ...),\ row_{out},\ row_{handled},\ [S\ N_1]]\!] \\ \text{fresh-var}[\![N_1,\ t_v,\ N_2]\!] \\ \text{fresh-var}[\![N_2,\ t_r,\ N_3]\!] \\ (x_v\ t_v\ (x_r\ (t_r\ \text{->}\ row_{out}\ t_{ret})\ \Gamma))\ |\ [S\ N_3]\ \vdash\ e:t_h\ !\ row_h\ |\ SN_2 \\ SN_2\ row_{out}\ \sim\ row_h\ SN_3 \\ SN_3\ t_{ret}\ \sim\ t_h\ SN_4 \\ \text{trace}[\![(\text{"<<< inferred handlers <<<}\backslash n{\sim}s\backslash n{\sim}s\backslash n\text{"} \\ \text{apply}[\![\text{subst}[\![SN_4]\!],\ t_{ret}]\!]\ \text{apply}[\![\text{subst}[\![SN_4]\!],\ row_{out}]\!])]\!] \end{array}}{\text{infer-handlers}[\![\Gamma,\ SN_1,\ t_{ret},\ ((op\ (x_v\ x_r\ e))\ h\ ...),\ row_{out},\ (op\ (t_v\ \text{=>}\ t_r)\ row_{handled}),\ SN_4]\!]}$$

Figure 2.3: Handlers type inference

it's argument. Accordingly the environment gets extended with functional type $t_1 \to row_1 t_2$ to check the body of the function, and afterwards the result type of body t gets unified with the result type of function t_2, same with effect row. Whole function, as it is a value, is returned with fresh effect row. The application requires expression at function position to be of functional type and parameter type to unify with argument type. All effect rows (from evaluation of function value, argument value and function body) must unify as well. The checking of primitive operation gets deferred to auxiliary judgment, which checks arity and argument types, returning result type and usually fresh effect row. Conditional expression requires the condition expression to be of type *Bool* and types of two branches to unify. As usual all effect rows must also unify. Operation invocation requires the effect row to contain operation *op* whose input type is the inferred type for e and output type being fresh. Operation lifting prepends fresh *op* to the effect row of e. Finally, to check handle expression we first infer the type of enclosed expression e, then in environment extended with e's type t_1 we infer return expression's type $t_r et$. Helper judgment *infer-handlers* returns the result effect row of handlers row_{out} and row marking handled effects $row_{handled}$ whose tail is the same as result's. By unifying result row with return row and handled row with row_1 we ensure that effects which may occur during handling of operations, at return and leftovers from the inner expression are all accounted for.

2.1.3 Effect handlers

List of effect handlers hs is processed right-to-left by judgment

$$infer\text{-}handlers[\Gamma,\ SN_{in},\ t_{ret},\ hs,\ row_{out},\ row_{handled},\ SN_{out}]$$

The t_{ret} is the type of return clause, row_{out} is the combined row of effects which may occur in any handler and $row_{handled}$ is the row of handled operations, with

appropriate types. The base case of empty list initializes both rows with the same type variable, this way $row_{handled}$ returned by $infer\text{-}handlers$ judgment will first have all handled operations and it's tail will be the same as row_{out}. The inductive case first calculates row_{out} and $row_{handled}$ for the tail of the list. Then two new fresh variables are created: t_v for the type of value passed to handler and t_r for type of resumption parameter. Resumption's result type is the result type of return clause as it should eventually evaluate to a value, which will be transformed by that clause. It's effect row is the same as the result row of whole handle expression, this means that any subsequent uses of operations will be handled by this handler. With environment extended with t_v for first parameter - an argument to handler expression and $t_r \rightarrow row_{out} \ t_{ret}$ for second parameter - the resumption, we check the handler expression e. We then unify the effects of evaluating e with all effects of handlers, and result type t_h with return type t_{ret}. At last we extend handled row with current operation $(op(t_v \Rightarrow t_r))$

2.2 Dynamic semantics

The dynamic semantics is defined in the format of contextual reduction semantics. With syntax of execution contexts in Figure 2.1 we see that they describe computation in call-by-value order, with expressions in applications (both normal and primitive) being evaluated left to right. All rules perform reduction in some context E leaving it unchanged. First rule describes standard β reduction via substitution of argument for variable in function body, while the second β reduction of recursive function, where first we substitute the function for function variable and then substitute the argument. Next rule deals with built-in primitive operations, delegating to auxiliary function $prim\text{-}apply$ which pattern matches on $prim$ and performs appropriate operation. Next two rules perform choice of correct branch in conditional expression depending on value of condition. The rule $lift\text{-}compat$ returns value from lift expression, leaving it unchanged. The rule $handle\text{-}return$ handles the case when inner expression of handle expression evaluates to a value, which means we have to evaluate return clause - by substituting the result value for x, and plugging this expression into evaluation context E. The last rule $handle\text{-}op$ describes the behavior when an expression calls some operation. Before we can define it's behavior we must define notion of $free$ness. The judgment $free[op, E, n]$ asserts that operation op is n free in evaluation context E, which means that it will be handled by $n + 1$-st handler $outside$ the context E [2]. With this definition in mind, we can see that to handle an operation we must find 0-$free$ inner context E_2 which is directly surrounded by handle expression which has a case for op. Then we substitute value v for first variable of operation handler and the inner context E_2 surrounded with the very same handler (the continuation delimited by the handler) closed in a lambda for the second argument. This way the operation handler can resume the evaluation of expression which invoked the handled operation.

$$E[((\lambda\ x\ e)\ v)] \longrightarrow E[\text{substitute}[\![e,\ x,\ v]\!]] \tag{β-λ}$$

$$E[(((\text{rec}\ x_f\ x_a\ e)\ v)] \longrightarrow E[\text{substitute}[\![\text{substitute}[\![e,\ x_f,\ (\text{rec}\ x_f\ x_a\ e)]\!],\ x_a,\ v]\!]] \tag{β-rec}$$

$$E[(prim\ v_1\ v_2)] \longrightarrow E[\text{prim-apply}[\![prim,\ v_1,\ v_2]\!]] \tag{prim-op}$$

$$E[(\text{if true}\ e_1\ e_2)] \longrightarrow E[e_1] \tag{if-true}$$

$$E[(\text{if false}\ e_1\ e_2)] \longrightarrow E[e_2] \tag{if-false}$$

$$E[(\text{lift}\ op\ v)] \longrightarrow E[v] \tag{lift-compat}$$

$$E[(\text{handle}\ v\ hs\ (\text{return}\ x\ e))] \longrightarrow E[\text{substitute}[\![e,\ x,\ v]\!]] \tag{handle-return}$$

$$E_1[(\text{handle}\ E_2[(op\ v)]\ hs\ ret)] \longrightarrow E_1[\text{substitute}[\![\text{substitute}[\![e,\ x_1,\ v]\!], \tag{handle-op}$$
$$x_2,\ (\lambda\ \text{v:z}\ (\text{handle}\ E_2[\text{v:z}]\ hs\ ret))]\!]]$$
$$\text{where free}[\![op,\ E_2,\ 0]\!],\ \text{get-handler}[\![op,\ hs,\ (x_1\ x_2\ e)]\!],\ \text{v:z fresh}$$

Figure 2.4: Reduction relation

2.3   Abstract machine

Chapter 3

Implementation

3.1 `PLT Redex`

3.2 Typing relation

3.3 Unification

3.4 Reduction relation

3.5 Automatic testing

Chapter 4

The `Racket` environment

4.1   Front-end

4.2   Back-end

Chapter 5

User's manual

Bibliography

[1] Andrej Bauer and Matija Pretnar. ?Programming with Algebraic Effects and Handlers? In: *CoRR* abs/1203.1539 (2012). arXiv: 1203.1539. URL: http://arxiv.org/abs/1203.1539.

[2] Dariusz Biernacki et al. ?Handle with Care: Relational Interpretation of Algebraic Effects and Handlers? In: *Proc. ACM Program. Lang.* 2.POPL (Dec. 2017), 8:1–8:30. ISSN: 2475-1421. DOI: 10.1145/3158096. URL: http://doi.acm.org/10.1145/3158096.

[3] Matthias Felleisen, Robert Bruce Findler, and Matthew Flatt. *Semantics Engineering with PLT Redex*. 1st. The MIT Press, 2009. ISBN: 0262062755, 9780262062756.

[4] Daniel Hillerström and Sam Lindley. ?Liberating Effects with Rows and Handlers? In: *Proceedings of the 1st International Workshop on Type-Driven Development*. TyDe 2016. Nara, Japan: ACM, 2016, pp. 15–27. ISBN: 978-1-4503-4435-7. DOI: 10.1145/2976022.2976033. URL: http://doi.acm.org/10.1145/2976022.2976033.

[5] Oleg Kiselyov and KC Sivaramakrishnan. ?Eff Directly in OCaml? In: *arXiv e-prints*, arXiv:1812.11664 (Dec. 2018), arXiv:1812.11664. arXiv: 1812.11664 [cs.PL].

[6] Daan Leijen. ?Extensible records with scoped labels? In: *Proceedings of the 2005 Symposium on Trends in Functional Programming (TFP'05), Tallin, Estonia*. Sept. 2005. URL: https://www.microsoft.com/en-us/research/publication/extensible-records-with-scoped-labels/.

[7] Daan Leijen. ?Koka: Programming with Row Polymorphic Effect Types? In: *Mathematically Structured Functional Programming 2014*. EPTCS, Mar. 2014. URL: https://www.microsoft.com/en-us/research/publication/koka-programming-with-row-polymorphic-effect-types-2/.

[8] Gordon Plotkin and John Power. ?Algebraic Operations and Generic Effects? In: *Applied Categorical Structures* 11.1 (Feb. 2003), pp. 69–94. ISSN: 1572-9095. DOI: 10.1023/A:1023064908962. URL: https://doi.org/10.1023/A:1023064908962.

[9]   Gordon Plotkin and Matija Pretnar. **?**Handling Algebraic Effects**?** In: *Logical Methods in Computer Science* 9.4 (Dec. 2013). Ed. by Andrzej Tarlecki. ISSN: 1860-5974. DOI: `10.2168/lmcs-9(4:23)2013`. URL: `http://dx.doi.org/10.2168/LMCS-9(4:23)2013`.